

Intro To Socket Programming

By: Matthew Johnson

Sockets

- Endpoint of two-way communication between two programs
- UDP & TCP Sockets
 - UNIX
 - C
- Use the sockets API
 - `sys/socket.h`

Client & Server

- Client
 - Needs resources from a server
 - Initiates a connection with a server
 - To create connection, client needs server's IP address and port number
- Server
 - Contains resources needed by clients
 - Waits for connections from clients on a specified port
 - Once a client initiates a connection, the server will have the client's IP address and port number

Socket Data Structures

- `int socket(int domain, int type, int protocol)`
 - `socket()` creates a socket to be used by the program
 - Returns socket descriptor of type `int` on success and `-1` on failure
 - `domain` is the communication domain
 - For the labs we will use the constant `PF_INET`
 - `type` is the type of socket to be created
 - For lab 2 we will use `SOCK_STREAM` for TCP communication
 - For lab 3 we will use `SOCK_DGRAM` for UDP communication
 - `protocol` is the protocol used by the socket
 - For the labs we will use `0`
 - Example: `int sock = socket(PF_INET, SOCK_DGRAM, 0)`

Socket Data Structures Cont.

- Specifying Server Address for clients
- struct sockaddr_in addr
 - addr.sin_family
 - Communication domain
 - Set to PF_INET
 - addr.sin_port
 - Server port to send message to
 - Set to htons(portNum)
 - htons(int port) takes an integer and returns a network byte order short
 - addr.sin_addr.s_addr
 - Server IP address to send message to
 - Set to inet_addr(ipDest)
 - inet_addr(char *) takes a string and returns an IP address

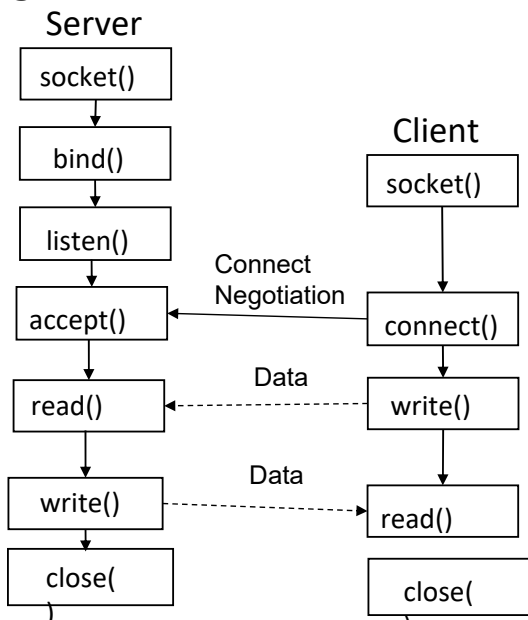
Socket Data Structures Cont.

- Specifying Client Address for servers
- struct sockaddr_in addr
 - addr.sin_family
 - Communication domain
 - For lab 3 PF_INET
 - addr.sin_port
 - Server port to receive messages into
 - Set to htons(portNum)
 - htons(int port) takes an integer and returns a network byte order short
 - addr.sin_addr.s_addr
 - Set to dummy IP address
 - Set to INADDR_ANY

TCP Sockets

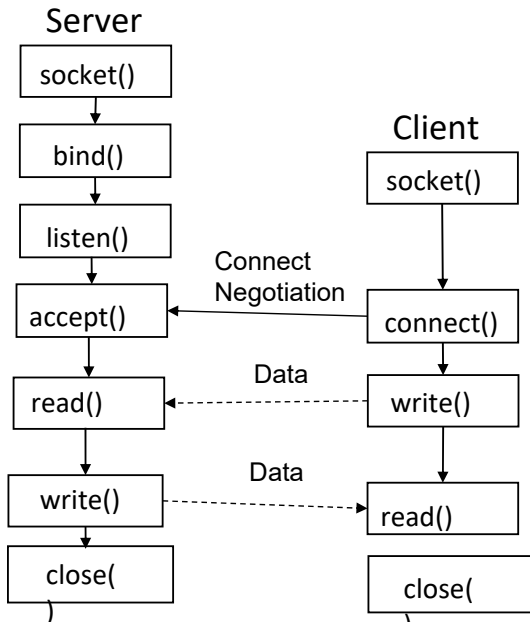
- Used in Lab 2
- Set up communication channel between server and client
- Server listens to a port for incoming communication requests
- Client connects to a server port to begin communication

TCP Server



1. Server creates socket using `socket()`
 - a. `socket(int domain, int type, int protocol)`

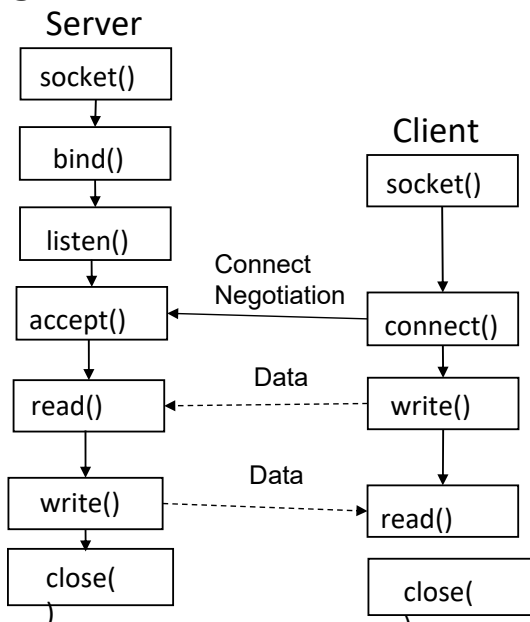
TCP Client



1. Client creates socket using `socket()`

a. `socket(int domain, int type, int protocol)`

TCP Server



2. Server binds to port using `bind()`

a. `bind(int sockfd, const struct sockaddr *addr, socklen_t addrlen)`

b. Allows server to bind to a port

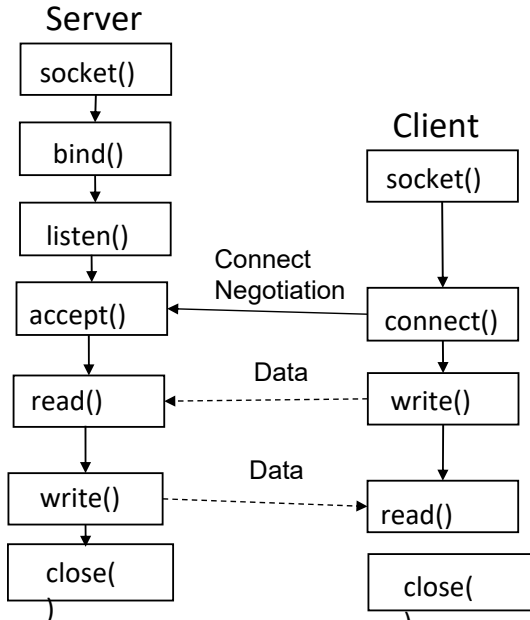
c. `sockfd` = socket created in step 1

d. `addr` = struct containing server information

e. `addrlen` = `sizeof(serverAddr)`;

f. returns 0 on success, -1 on failure

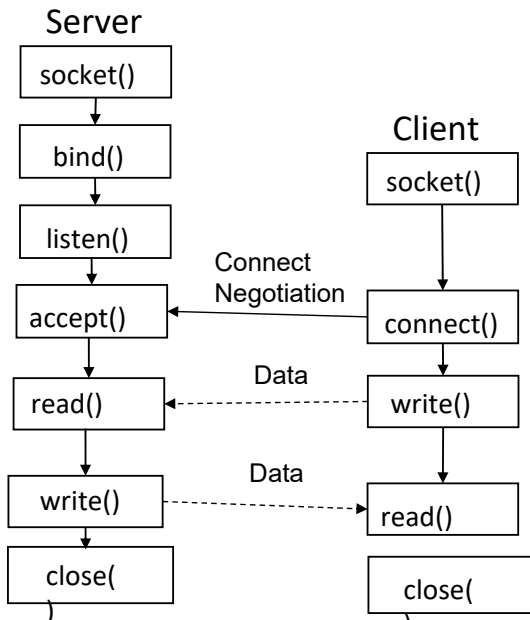
TCP Server



3. Server listens to port using listen()

- `listen(int sockfd, int backlog)`
- Allows server to listen to a port
- `sockfd` = socket created in step 1
- `backlog` = number of allowed pending connections
- returns 0 on success, -1 on failure

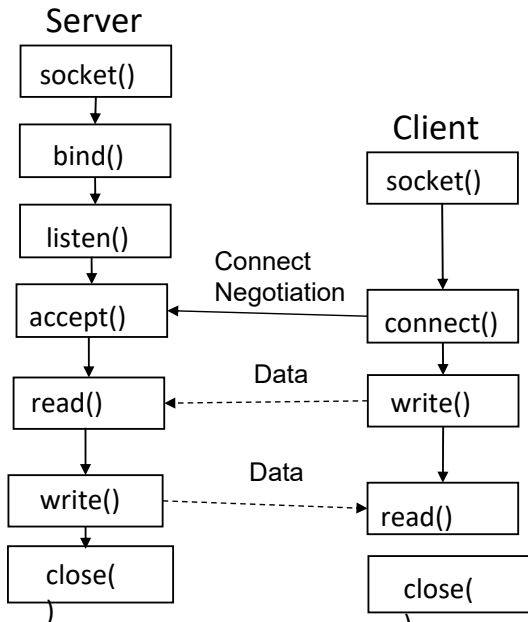
TCP Server



4. Server accepts an incoming connection using accept()

- `accept(int sockfd, struct sockaddr *cliaddr, int *addrlen)`
- Blocking while it waits for a connection
- Accepts an incoming connection request
- `sockfd` = socket created in step 1
- `cliaddr` = pointer to `sockaddr` struct where client's address will be stored
- `addrlen` = pointer to integer where length of client address will be stored
- returns new socket on success, -1 on failure


TCP Client



2. Client attempts to connect to server port using `connect()`

- `connect(int sockfd, struct sockaddr *seraddr, socklen_t addrlen);`
- Attempts to make connection between client socket and server socket using TCP 3-way handshake
- `sockfd` = socket created in step 1
- `seraddr` = pointer to `sockaddr` struct where server's address and port are specified
- `addrlen` = length of server address struct
- returns 0 on success, -1 on failure

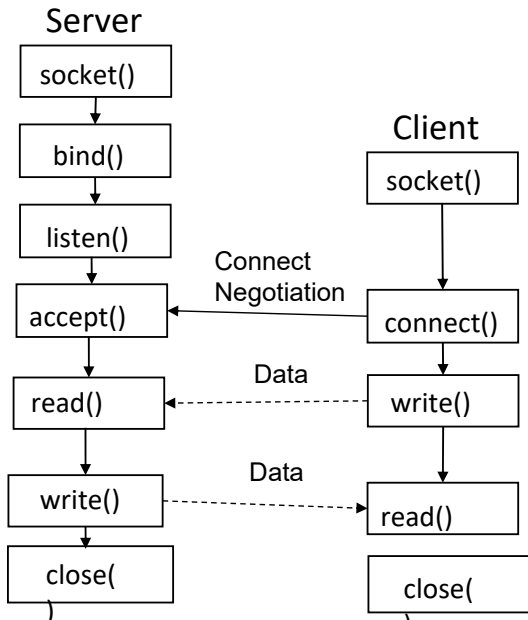
TCP Client



3. Client writes data over the socket using `write()`

- `write(int fd, char *buf, int num)`
- `fd` = socket created in step 1
- `buf` = char buffer that holds data to send
- `num` = number of bytes to send
- returns number of bytes sent on success, -1 on failure


TCP Server



5. Server reads data sent by client using `read()`

- `read(int fd, char *buf, int max)`
- `fd` = socket created by `accept()` in step 4
- `buf` = char buffer that will hold received data
- `max` = maximum number of bytes that can be read
- returns number of bytes read on success, -1 on failure


TCP Server



6. Server writes data over the socket using `write()`

- `write(int fd, char *buf, int num)`
- `fd` = socket created in `accept()` in step 4
- `buf` = char buffer that holds data to send
- `num` = number of bytes to send
- returns number of bytes sent on success, -1 on failure

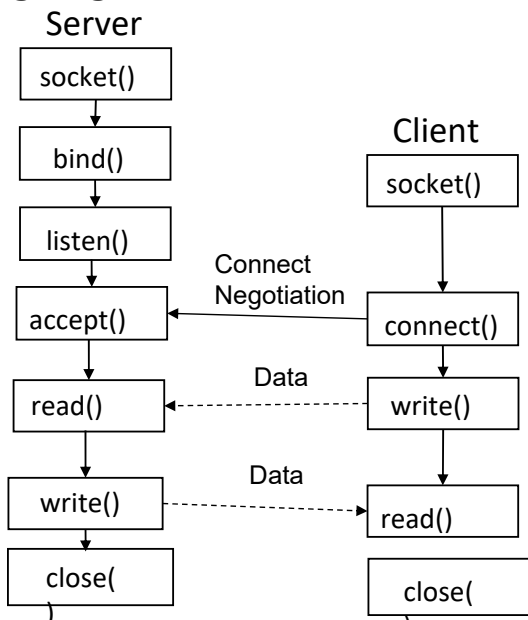
TCP Client



4. Client reads data sent by server using `read()`

- `read(int fd, char *buf, int max)`
- `fd` = socket created by step 1
- `buf` = char buffer that will hold received data
- `max` = maximum number of bytes that can be read
- returns number of bytes read on success, -1 on failure


TCP Client



5. Client closes socket using `close()`

- `close(int fd)`
- Closes connection
- `fd` = socket created by step 1
- returns number of 0 on success, -1 on failure

TCP Server




7. Server closes socket using `close()`

- `close(int fd)`
- Closes connection
- `fd` = socket created by step 1
- returns number of 0 on success, -1 on failure

UDP Sockets


- Used in Lab 3
- Communication by datagram
- Server binds to socket and waits for incoming messages
- Client directly sends messages to a server's socket

UDP Server




1. Server creates socket using `socket()`
 - a. `socket(int domain, int type, int protocol)`

UDP Client



1. Client creates socket using `socket()`
 - a. `socket(int domain, int type, int protocol)`


UDP Server



2. Server binds to port using bind()

- `bind(int sockfd, const struct sockaddr *addr, socklen_t addrlen)`
- Allows server to bind to a port
- `sockfd` = socket created in step 1
- `addr` = struct containing server information
- `addrlen` = `sizeof(serverAddr)`;
- returns 0 on success, -1 on failure

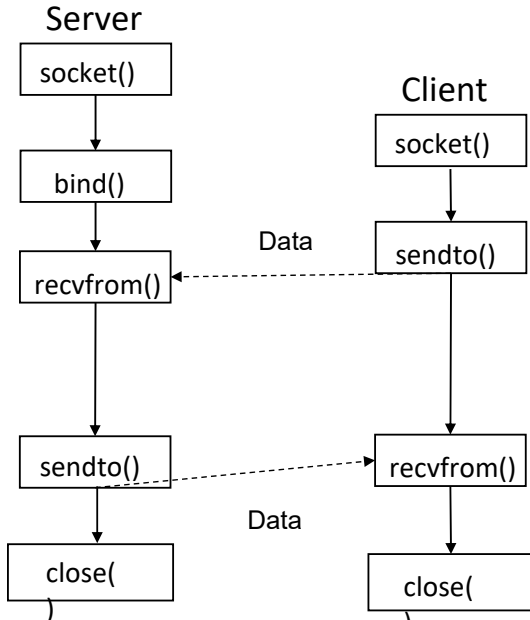
UDP Client



2. Client sends data to server

- `sendto(int sockfd, const void *buf, size_t len, int flags, const struct sockaddr *dest_addr, socklen_t addrlen)`
- `sockfd` = socket created in step 1
- `buf` = char buffer to send data from data into
- `len` = bytes of buffer to send
- `flags` = options to receive data (0 for lab 3)
- `dest_addr` = struct of server address
- `addrlen` = `sizeof(dest_addr)`
- returns number of bytes sent on success, -1 on failure

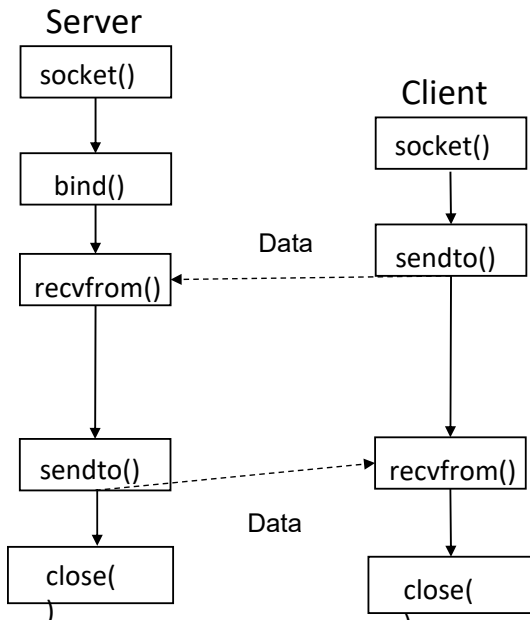
UDP Server



3. Server receives data from client(recvfrom is blocking)

- `recvfrom(int sockfd, void *buf, size_t len, int flags, struct sockaddr *src_addr, socklen_t *addrlen)`
- `sockfd` = socket created in step 1
- `buf` = char buffer to receive data into
- `len` = size of char buffer
- `flags` = options to receive data (0 for lab 3)
- `src_addr` = struct to be populated with source address (for lab 3 use blank `sockaddr_in` struct)
- `addrlen` = `sizeof(src_addr)`
- returns number of bytes received on success, -1 on failure

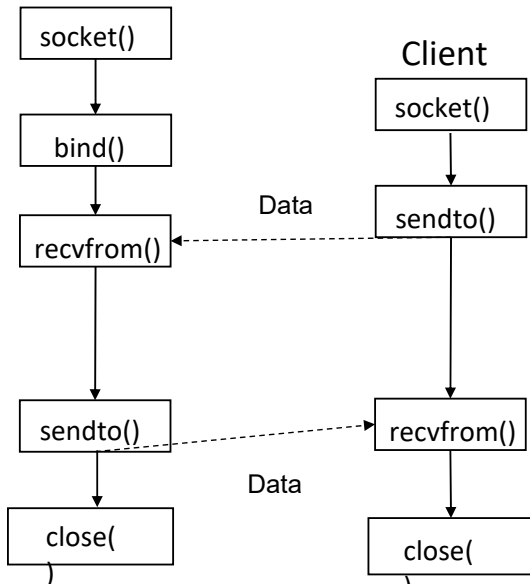
UDP Server



4. Server sends data to client

- `sendto(int sockfd, const void *buf, size_t len, int flags, const struct sockaddr *dest_addr, socklen_t addrlen)`
- `sockfd` = socket created in step 1
- `buf` = char buffer to send data from data into
- `len` = bytes of buffer to send
- `flags` = options to receive data (0 for lab 3)
- `dest_addr` = struct of client address populated in `recvfrom()`
- `addrlen` = `sizeof(dest_addr)`
- returns number of bytes sent on success, -1 on failure


UDP Client



3. Client receives data from server(recvfrom is blocking)

- recvfrom(int sockfd, void *buf, size_t len, int flags, struct sockaddr *src_addr, socklen_t *addrlen)
- sockfd = socket created in step 1
- buf = char buffer to receive data into
- len = size of char buffer
- flags = options to receive data (0 for lab 3)
- src_addr = server address struct
- returns number of bytes received on success, -1 on failure

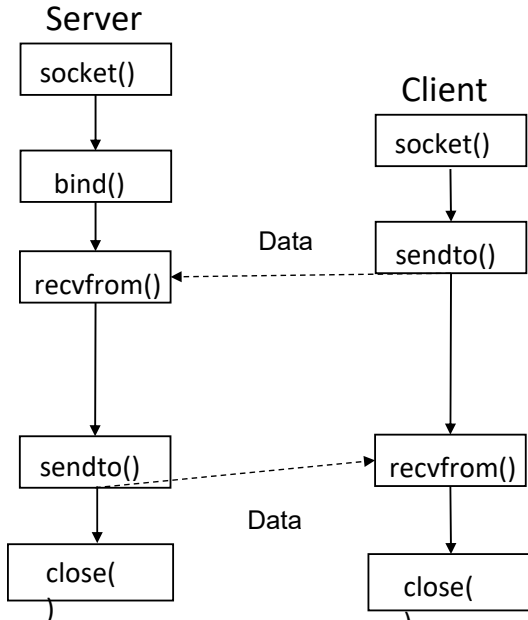
UDP Server



5. Server closes socket

- Closes socket connection
- close(int fd)
- fd = socket to close
- returns number of 0 on success, -1 on failure

UDP Client



4. Client closes socket

- Closes socket connection
- `close(int fd)`
- `fd` = socket to close
- returns number of 0 on success, -1 on failure

Helpful Tips

- A TCP server `accept()` returns a socket, this should be used for communication with a client
- Use error checking on socket api calls
 - `if(socket(PF_INET,SOCK_STREAM,0)<0) { perror("Socket Creation Failure: "); return -1;}`
- Make sure the correct number of bytes are being sent from the buffer
- Don't forget to format IP addresses and port numbers using `inet_addr()` and `htons()`
- Make sure all data structures and buffers are 0 filled using `memset` or other method
- Use correct values for length on socket api calls
- Make sure you are using the correct inputs to socket api calls
 - Pointers vs values